


성남-KAIST AI

[Deep Learning 실습]

TAs: Kiwon Lee and Bae Gwangtak

Advisor: Junmo Kim

KAIST SIIT Lab



[Contact Info.]


Kiwon Lee: kaiser5072@kaist.ac.kr

Bae Gwangtak: gwangtak.bae@kaist.ac.kr


Junmo Kim: junmo.kim@kaist.ac.kr

Tutorial Guidelines


- **Session 1. Supervised Learning (Image Classification)**
 - Convolutional Neural Network (DIY network)



- **Session 2. Unsupervised Learning (Image Generation)**
 - Generative Adversarial Network (GAN)



What is Tensorflow?



- A python library
- Google (Brain)
- Open-source
- Library for numerical computation using data flow graphs
- CPU and GPU

<https://www.tensorflow.org/>



Install Requirements

- **CPU-only**


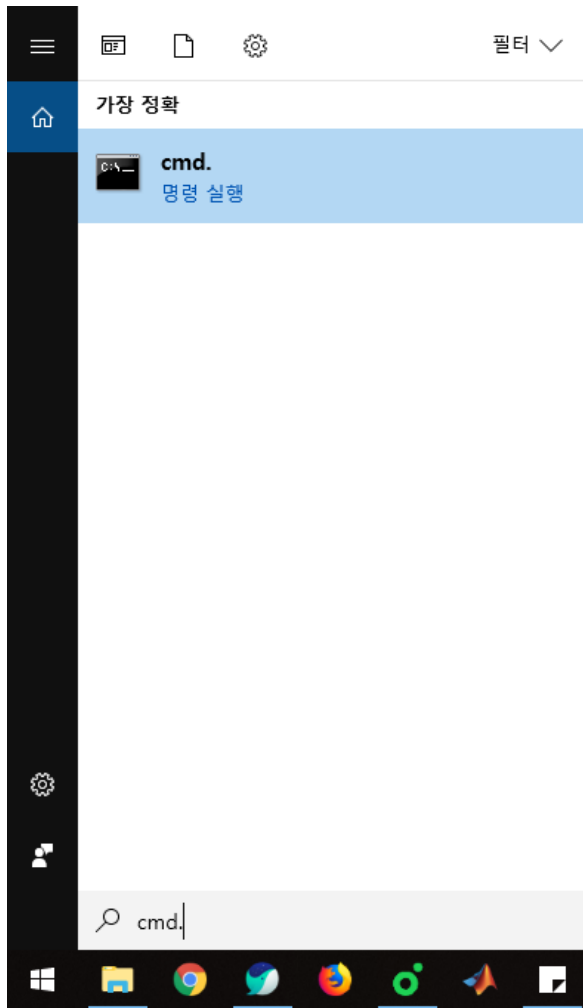
- Ubuntu 16.04 or later (64 bit)
- macOS 10.12.6 (Sierra) or later (64 bit)
- Windows 7 or later (64 bit) (Python 3 only)

- **GPU support**

- NVIDIA GPU card withi CUDA Compute Capability 3.5 or higher
<https://developer.nvidia.com/cuda-gpus>
- GPU drivers, CUDA Toolkit, cuDNN...

CPU-only Tensorflow (I)

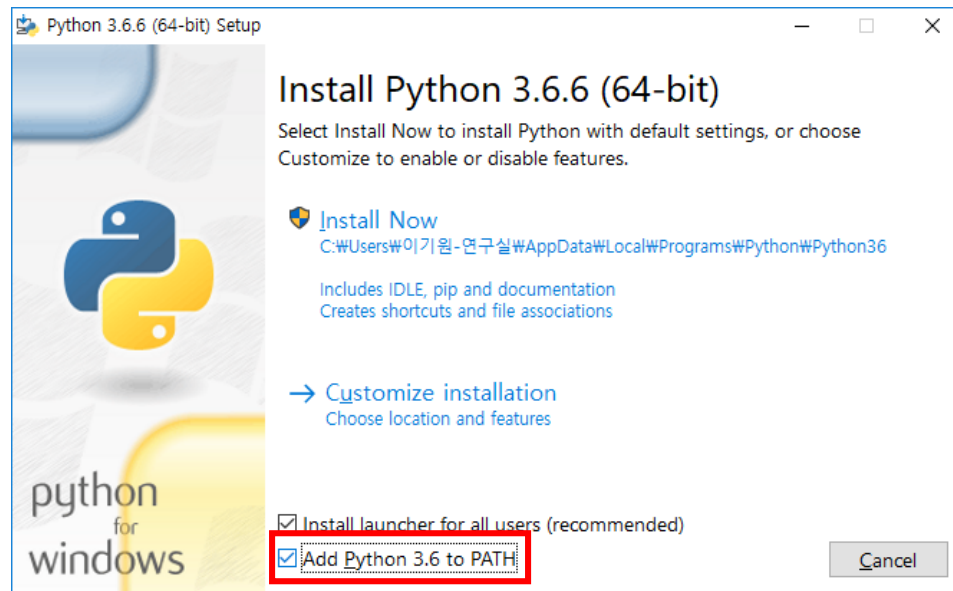
- Click windows + s button and execute terminal



```
python3 --version  
pip3 --version  
virtualenv --version
```

CPU-only Tensorflow (II)

- Update build tools
 1. Go to the [Visual Studio downloads](#)
 2. Select *Redistributables and Build Tools*
 3. Download and install the Microsoft Visual C++ 2015 Redistributable Update 3.
- Install [python 3](#)



CPU-only Tensorflow (III)

- Install virtualenv

```
pip3 install -U pip virtualenv
```

- Install Tensorflow

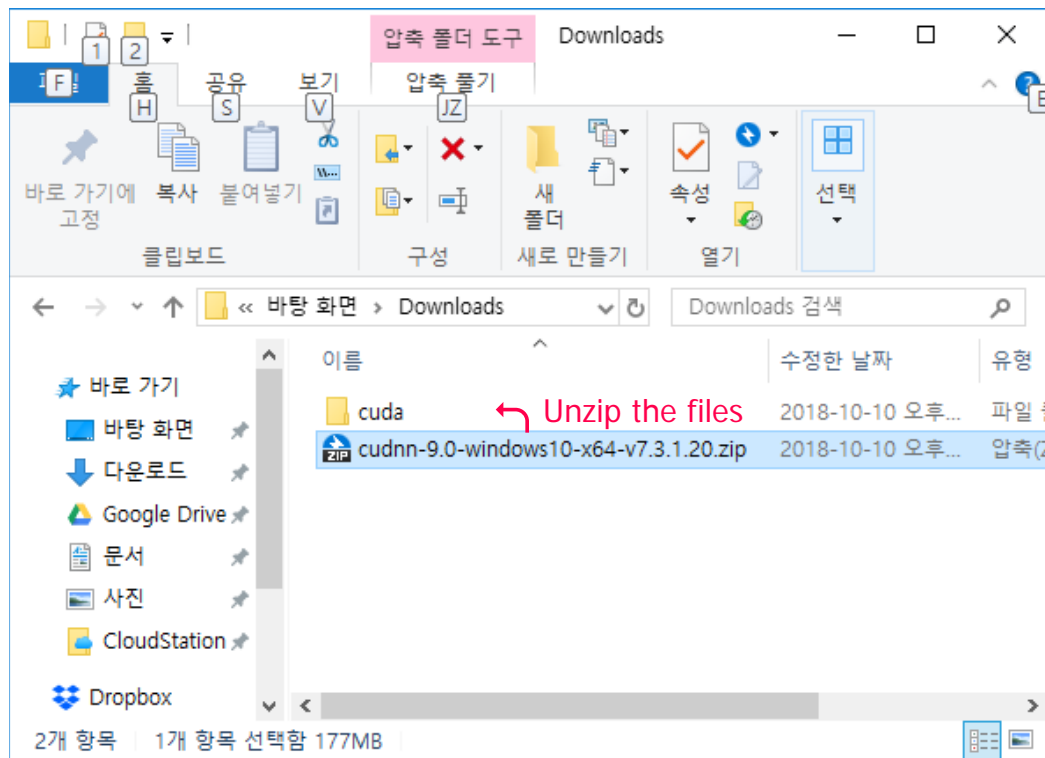
```
pip3 install -user -upgrade tensorflow  
python3 -c "import tensorflow as tf; print(tf.__version__)"
```

GPU support Tensorflow (I)

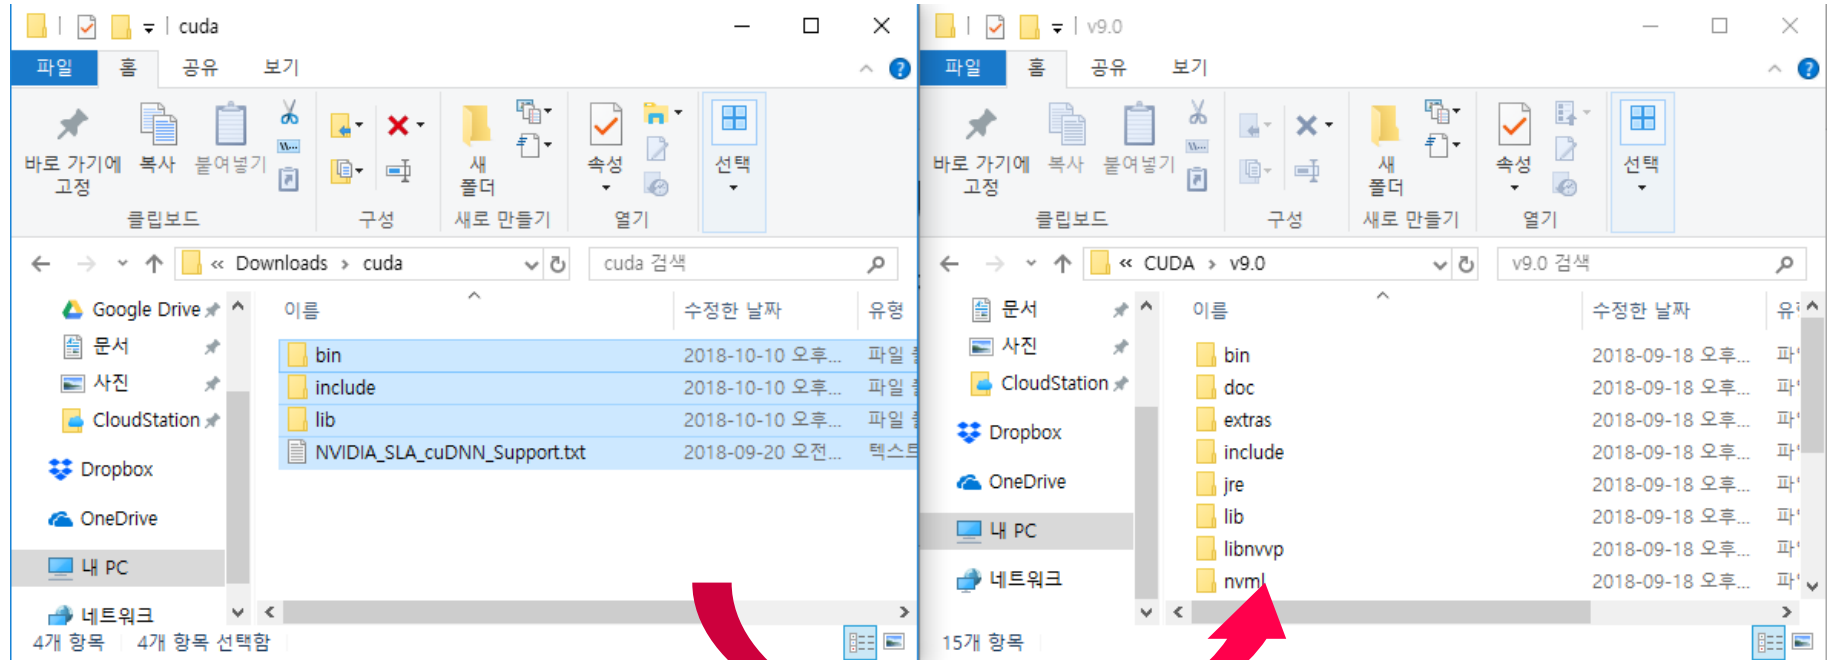
- Requirements

<https://www.tensorflow.org/install/gpu>

- [CUDA Toolkit](#) – Tensorflow supports CUDA 9.0
- [cuDNN SDK](#) – (≥ 7.2)



GPU support Tensorflow (II)



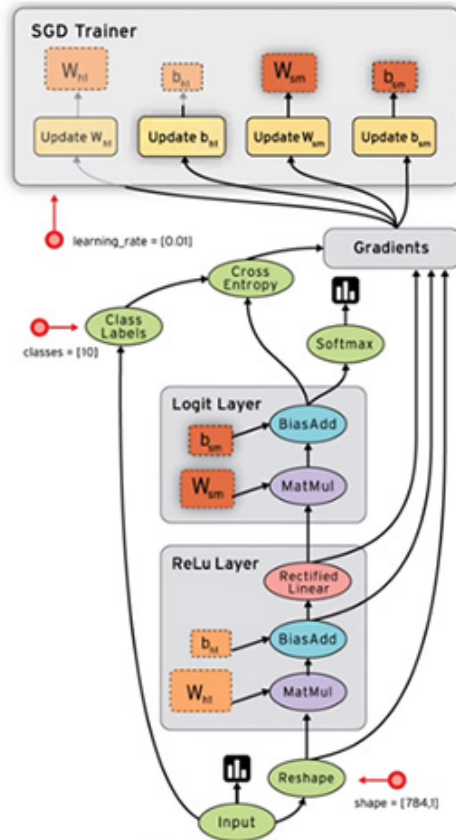
Copy the files to the folder where cuda is installed

GPU support Tensorflow (III)

- Install Tensorflow

```
pip3 install --user --upgrade tensorflow-gpu  
python3 -c "import tensorflow as tf; print(tf.__version__)"
```

Principle



[Tensorflow workflow]

1. Draw your graph
2. Feed data
3. ... and optimize

[Graph]

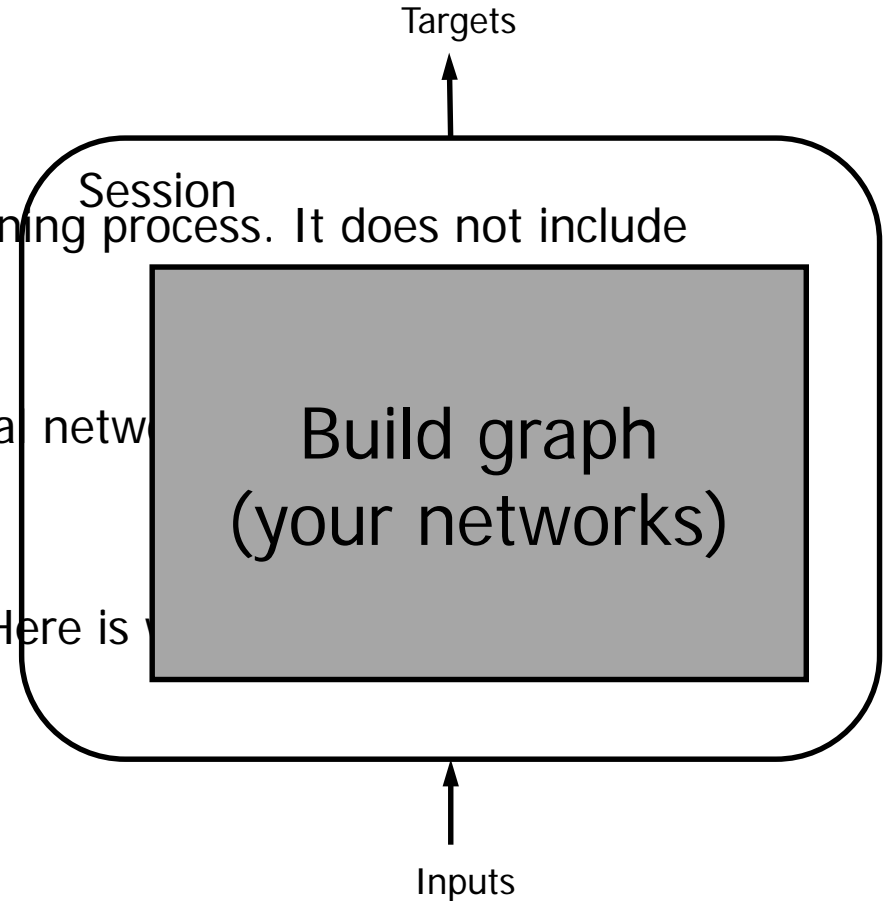
- Placeholders: gates where we introduce examples.
- Model: makes predictions. Set of variables and operations.
- Cost function: function that computes the model error.
- Optimizer: algorithm that optimizes the variables.

Example Code

- `import tensorflow as tf`
- `mnist = tf.contrib.learn.datasets.load_dataset()`
- `model.build_graph()`
- `sess = tf.Session()`
- `sess.run(tf.global_variables_initializer())`
- `for t in range(max_iters):`
 - `sess.run(model.cost, model.train_op)`

Graph, Data and Session

- Graph
 - Layout of the prediction and learning process. It does not include data.
- Data
 - Examples that will train the neural network inputs and targets.
- Session
 - Where everything takes places. Here is where you build your graph with data.



Constant, variable, and placeholder

- Difference between constant/variable and placeholder.
- Placeholder: Feeding data to your graph during session.
- Variable: learnable/not learnable
- Constant: a fixed variable

[Placeholder]

```
images = tf.placeholder(tf.float32, shape=(hps.batch_size, 28, 28, 1), name='images')
```

[Variable]

```
w = tf.get_variable('weight/DW', [x.get_shape()[1], out_dim],  
                    initializer=tf.uniform_unit_scaling_initializer(factor=1.0),  
                    trainable=True)
```

[Constant]

```
Init = tf.ones(...) / tf.random_normal(...)
```

Operations

- Tensorflow math ops are pretty standard, quite similar to Numpy.
- Visit https://www.tensorflow.org/api_docs/cc/group/math-ops

Category	Examples
Element-wise mathematical operations	Add, Sub, Mul, Div, Exp, Log, Greater, Less, Equal, ...
Array operations	Concat, Slice, Split, Constant, Rank, Shape, Shuffle, ...
Matrix operations	MatMul, MatrixInverse, MatrixDeterminant, ...
Stateful operations	Variable, Assign, AssignAdd, ...
Neural network building blocks	SoftMax, Sigmoid, ReLU, Convolution2D, MaxPool, ...
Checkpointing operations	Save, Restore
Queue and synchronization operations	Enqueue, Dequeue, MutexAcquire, MutexRelease, ...
Control flow operations	Merge, Switch, Enter, Leave, NextIteration

```
tf.add(a, b) # >> [5 8]
tf.add_n([a, b, b]) # >> [7 10]. Equivalent to a + b + b
tf.mul(a, b) # >> [6 12] because mul is element wise
tf.matmul(a, b) # >> ValueError
tf.matmul(tf.reshape(a, shape=[1, 2]), tf.reshape(b, shape=[2, 1])) # >> [[18]]
tf.div(a, b) # >> [1 3]
tf.mod(a, b) # >> [1 0]
```

Feeds and Fetches

- Feeds
 - Designating specific operations to be 'feed' by using `tf.placeholder()`.
- Fetches
 - Execute the graph with a `run()` call on the Session object and pass in the tensors to retrieve.
 - All the ops needed to produce the values of the requested tensors are run once.

```
input1 = tf.placeholder(tf.float32)
input2 = tf.placeholder(tf.float32)
input3 = tf.placeholder(tf.float32)
intermed = tf.add(input2, input3)
mul = tf.mul(input1, intermed)
```

With `tf.Session()` as `sess`:

```
result = sess.run([output],
                   feed_dict={input1: 3.0,
                               input2: 2.0,
                               input3: 5.0})
print(result)
```

```
input1 = tf.constant([3.0])
input2 = tf.constant([2.0])
```

```
input3 = tf.constant([5.0])
intermed = tf.add(input2, input3)
mul = tf.mul(input1, intermed)
```

With `tf.Session()` as `sess`:

```
result = sess.run([mul, intermed])
print(result)
```


Coding by Examples

Image Classification

- Task: image classification
- Dataset: MNIST (28x28)
- Network: Simple CNN (VGG-like)



1. Prepare your data
2. Build graph
3. Make session & Initialize
4. Run session

1. Prepare your data

- Load MNIST data

```
from tensorflow.examples.tutorials.mnist import input_data  
mnist = input_data.read_data_sets('./data/', one_hot=True)
```

- Split train/val/test (image, label) pairs

```
train_images = mnist.train.images  
train_labels = mnist.train.labels  
train_images = train_images.reshape([-1, 28, 28, 1]) / 255.  
train_images = train_images[0:1000]
```

```
val_images = mnist.validation.images  
val_labels = mnist.validation.labels  
val_images = val_images.reshape([-1, 28, 28, 1]) / 255.
```

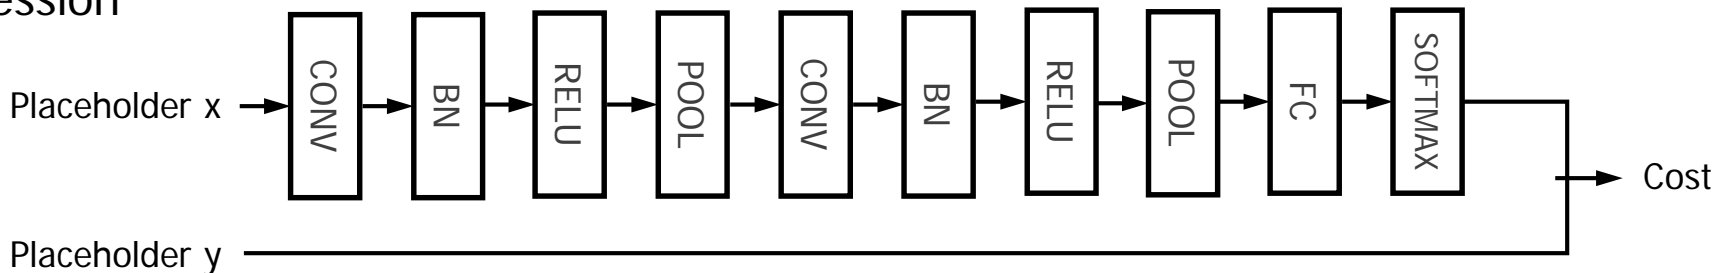
2. Build graph

- From inputs to targets.
- Data:

```
x = tf.placeholder(tf.float32, shape=[hps.batch_size, 28, 28, 1])  
y = tf.placeholder(tf.float32, shape=[hps.batch_size, 10])
```

- Network:
 - (CONV-BN-RELU) – POOL – (CONV-BN-RELU) – POOL – FC - SOFTMAX

Session



2. Build graph - Convolution

- Network:
 - (CONV-BN-RELU) – POOL – (CONV-BN-RELU) – POOL – FC - SOFTMAX

```
def conv(x, filter_size, in_filters, out_filters, strides=[1, 1, 1, 1]):  
    w = tf.get_variable(  
        'weight/DW', [filter_size, filter_size, in_filters, out_filters],  
        tf.float32, initializer=tf.uniform_unit_scaling_initializer(factor=2.0))  
    return tf.nn.conv2d(x, w, strides=[1, 1, 1, 1], padding='SAME')
```

Usage: `y = conv(x, 3, in_filter, out_filter, stride)`

2. Build graph – Batch normalization

- Network:
 - (CONV-BN-RELU) – POOL – (CONV-BN-RELU) – POOL – FC - SOFTMAX

```
def batch_norm(x):  
    params_shape = [x.get_shape()[-1]]  
  
    beta = tf.get_variable('beta', params_shape, tf.float32,  
        initializer=tf.constant_initializer(0.0, tf.float32))  
    gamma = tf.get_variable('gamma', params_shape, tf.float32,  
        initializer=tf.constant_initializer(1.0, tf.float32))  
    mean, variance = tf.nn.moments(x, [0, 1, 2])  
  
    return tf.nn.batch_normalization(x, mean, variance, beta, gamma)
```

Usage: $y = \text{batch_norm}(x)$

2. Build graph – ReLU

- Network:
 - (CONV-BN-RELU) – POOL – (CONV-BN-RELU) – POOL – FC - SOFTMAX

```
def relu(x, leakiness=0.0):  
    return tf.maximum(x, leakiness*x)
```

Usage: $y = \text{relu}(x, \text{leakiness})$

2. Build graph – Pooling

- Network:
 - (CONV-BN-RELU) – POOL – (CONV-BN-RELU) – POOL – FC - SOFTMAX

```
def max_pool(x):  
    return tf.nn.max_pool(x, ksize=[1, 2, 2, 1],  
                           strides=[1, 2, 2, 1], padding='SAME')  
  
def avg_pool(x):  
    return tf.nn.avg_pool(x, ksize=[1, 2, 2, 1],  
                           strides=[1, 2, 2, 1], padding='SAME')
```

Usage: $y = \text{max_pool}(x)$

2. Build graph – Fully connected

- Network:
 - (CONV-BN-RELU) – POOL – (CONV-BN-RELU) – POOL – FC - SOFTMAX

```
def fully_connected(self, name, x, out_dim):  
    x = tf.reshape(x, [self.hps.batch_size, -1])  
    w = tf.get_variable(  
        'weight/DW', [x.get_shape()[1], out_dim],  
        initializer=tf.uniform_unit_scaling_initializer(factor=1.0))  
    b = tf.get_variable('bias', [out_dim],  
        initializer=tf.constant_initializer())  
    return tf.nn.xw_plus_b(x, w, b)
```

Usage: `y = fully_connected(x, num_classes)`

2. Build graph – Cost function & Optimizer

- Network:
 - (CONV-BN-RELU) – POOL – (CONV-BN-RELU) – POOL – FC - SOFTMAX

Cost function
(Cross entropy)

```
cent = tf.nn.softmax_cross_entropy_with_logits(logits=logits, labels=self.labels)
self.cost = tf.reduce_mean(cent, name='cent')
```

Optimizer
(SGD+momentum)


```
def build_train_op():
    optimizer = tf.train.MomentumOptimizer(self.lrn_rate, 0.9)
    train_vars = tf.trainable_variables()
    grads = tf.gradients(self.cost, train_vars)
    apply_op = optimizer.apply_gradients(zip(grads, train_vars),
                                         global_step=self.global_step,
                                         name='train_step')
    train_ops = [apply_op] + self._extra_train_ops
    self.train_op = tf.group(*train_ops)
```


Experiment

Image Classification

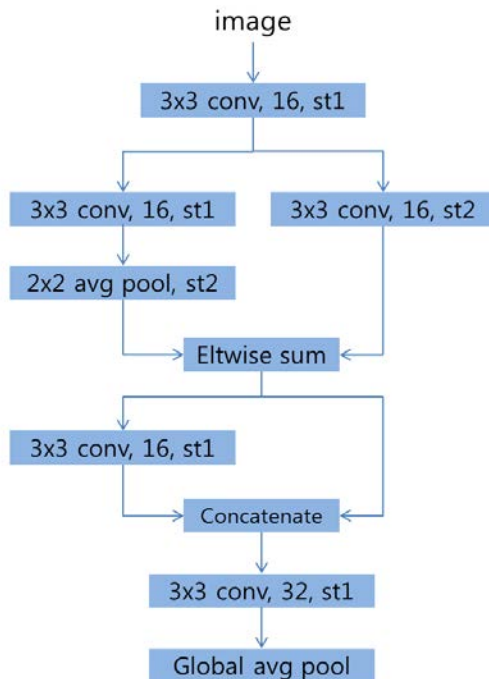
Task: Image Classification

- MNIST classification
- Provided code: VGG / ResNet architecture.



Task: Image Classification

- Project goal 1: (Simple Network)
 - Implement the following architecture.
- Project goal 2: (DIY Network)
 - Make your own architecture.



```
## BUILD GRAPH
## Option 1: resnet
## Option 2: vggnet
## Option 3: simple network (fill in the blank!)
## Option 4: DIY network (fill in the blank!)

def build_graph(self):
    ## GLOBAL ITERATION
    self.global_step = tf.contrib.framework.get_or_create_global_step()

    if self.mode == 'train':
        ## FEATURE EXTRACTION
        with tf.variable_scope('embed') as scope:
            #feats = self.resnet(self.images)
            #feats = self.vggnet(self.images)
            feats = self.simple_network(self.images)
            #feats = self.DIY_network(self.images)
```

```
##### FILL IN THE BLANK #####
##### BE CREATIVE! #####

def DIY_network(self, images):
    """ Your Code """

    img_feat = self._global_avg_pool(x)
    return img_feat
```